

# Synchronization

Process Coordination

# Race Condition

- A **race condition** is when multiple processes concurrently read and write to a shared memory location and the result depends on the order of execution
- Practical example: Account Holder 1 (Husband) and Holder 2 (Wife) withdraw ₹20000 cash simultaneously from different ATMs, when account had only ₹30000 to begin with.

# Race Condition in Programming

- Parallel execution of count++ and count--
- count++ implementation

```
register1 = count  
register1 = register1 + 1  
count = register1
```

- count-- implementation

```
register2 = count  
register2 = register2 - 1  
count = register2
```

# Execution order differences

# Initial value = 5, Final Result = 5

```
register1 = count  
register1 = register1 + 1  
count = register1
```

```
register2 = count  
register2 = register2 - 1  
count = register2
```

# Initial value = 5, Final Result = 4

```
register1 = count
```

```
register1 = register1 + 1
```

```
count = register1
```

```
register2 = count
```

```
register2 = register2 - 1
```

```
count = register2
```

# Initial value = 5, Final Result = 6

```
register1 = count
```

```
register2 = count  
register2 = register2 - 1  
count = register2
```

```
register1 = register1 + 1  
count = register1
```

# Solutions?

- Atomic Operations
- Critical Section Selection



# Atomics

- Group of operations is called **atomic** if all of them must be conducted as if it were a single operation
- Atomic operations **cannot** be interrupted by other operations.
- If due to any reason atomic operation set is not fully executed, entire atomic block is rolled back

# Atomics


```
atomic {  
    register1 = count  
    register1 = register1 + 1  
    count = register 1  
}
```

# Producer – Consumer Problem

A generic form of previously stated problem

Considered as a classical computer science problem, and significant in the history of multi-process operating systems

# Producer and Consumer



# Producer

```
while (true) {  
    nextProduced = produce();  
    while (counter == BUFFER_SIZE) {  
        // just wait when all slots are full  
    }  
    buffer[in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

# Consumer

```
while (true) {  
    while (counter == 0) {  
        // just wait when all slots are empty  
    }  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    consume(nextConsumed)  
}
```

# Critical Section Selection

# Shared Resources

- Resources ?
  - Computer components that a process needs to execute itself
    - memory space
    - filesystem access
    - I/O device access
- Shared Resource
  - When the same resource needs to be made available to **2 or more** processes
  - There is a probability of more than 1 process simultaneously accessing and modifying the quantity/quality of the resource



# Critical Section

- A **critical section** is part of program that accesses shared resource that cannot be concurrently accessed (without adverse effect)
- For example, during `count++` or `count--`, reading and writing the 'count' variable in memory are critical sections.
- Only **one instance** of a critical section must execute at a time on a shared resource
- Critical section bugs are the hardest to solve in OS

# Typical program with critical section

```
do {  
    entry section  
  
    critical section  
  
    exit section  
  
    remainder section  
  
} while (true);
```

# Critical Section Solutions

The solution **must enforce all three** required conditions -

- **MUTUAL EXCLUSION** – If process  $P_i$  is executing in its critical section, then no other processes can be executing in their critical sections
- **PROGRESS** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely
- **BOUNDED WAITING** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted

# Peterson's Solution

- Gary L Peterson; 1981
- Original solution: Works on only 2 processes
  - Space required: 2 x booleans, 1 x integer

# Initial condition

```
bool flag[2] = {false, false};  
int turn;
```

# Critical section selection

```
do {  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j) {  
        // wait  
    }  
    // critical section  
    flag[i] = false;  
    //remainder section  
} while (true)
```

# Bakery Algorithm

- Developed by Leslie Lamport
  - Scales Peterson's Solution to N processes
- Uses tokens at a bakery counter for analogy (hence the name)

# Variables required

- Given, **N** number of processes; we need the following –
- An integer array of length N (initially all 0)
  - `int number[N]`
- A Boolean array of length N (initially all false)
  - `boolean choosing[N]`



# The steps for critical section entry

- When process wants to enter critical section, indicate that it is about to get a choosing number
- It then receives a number. (number[i] for  $P_i$ )
- It indicates then that it has picked up a number
- Process with smallest choosing number enters critical section
- *If two processes have same token number, the one with lower index goes in.*
- *i.e., if `number[i] == number[j]`, compare i and j. Lower wins*

# Utilities and Assumptions

- `maxInArray()` finds largest number in a an array
- When comparing a,b with c,d we are check –
  - Is  $A < C$  ? If yes, then  $A, B < C, D$
  - If  $A == C$ , then is  $B < D$  ? If yes, then  $A, B < C, D$

# Code

```
do {  
    choosing[i] = true;  
    number[i] = maxInArray(number) + 1;  
    choosing[i] = false;  
    for (j = 0; j < n; j++) {  
        while (choosing[j]); //wait  
        while ((number[j] != 0) &&  
            (number[j], j < number[i], i)); //wait  
    }  
    //critical section  
    number[i] = 0;  
    //remainder section  
} while (true)
```

# Hardware based solutions

- Modern systems have hardware support for synchronizations
  - Uniprocessors – disables interrupts
  - Atomics – has atomic operation support
  - Test and Set – instruction support

# Uniprocessor

- Currently running code executes without preemption
- Too inefficient for multitasking systems
- Not scalable beyond application specific processors

# Test and Set (Hardware Atomic Instruction)

```
bool testAndSet (bool *target) {  
    bool origVal = *target;  
    *target = TRUE;  
    return origVal;  
}
```

# Critical Section Entry via TestAndSet

```
do {  
    while(testAndSet(&lock)); //wait  
  
    //critical section  
  
    lock = FALSE;  
  
    //remainder section  
} while (TRUE)
```

# Swap

```
void swap(int *a, int *b) {  
    int temp = *a;  
    *a = *b;  
    *b = temp;  
}
```




# Critical Section Entry via Swap

```
do {  
    key = TRUE;  
    while (key == TRUE)  
        swap(&lock, &key);  
  
    //critical section  
    lock = FALSE;  
    //remainder section  
} while (TRUE)
```

# Semaphores

- A hand signal using flags



# Semaphores

- ~~Magical integers that makes OS run 5x faster~~
- An integer that can be modified by ***wait()*** an ***signal()*** operations only
  - Helps achieve mutual exclusion

# Semaphore Operations (both atomic)

```
wait(S) {  
    while (S <= 0); //do nothing  
    S--;  
}
```

```
signal(S) {  
    S++;  
}
```

# Semaphores

- Two types –
  - counting (  $0 \rightarrow N$  )
  - binary (true / false)
- Binary Semaphores == Mutex Locks
  - They provide mutual exclusion support
- Counting semaphores control access to limited instance resource.

# Critical Section for N Processes

- Shared Data –

- semaphore mutex = 1;

- Process  $P_i$

```
do {  
    wait(mutex);  
    //critical section  
    signal(mutex);  
    //remainder section  
} while (TRUE);
```

# Producer Consumer – using Semaphores

- Shared Data –
  - `semaphore full, empty; //counting semaphores`
  - `semaphore mutex; //binary`
- Initial State
  - `full = 0`
  - `empty = N`
  - `mutex = 1`



# Producer Process

```
do {  
    // produce nextItem (remainder)  
    wait(empty);           // reserved empty slot first  
    wait(mutex);          // blocks the buffer for operation  
    // add nextItem to buffer (critical)  
    signal(mutex);        // free up buffer  
    signal(full);          // notify that one more slot is full  
} while (TRUE);
```

# Consumer Process

```
do {  
    wait(full);           // stake claim to an item first  
    wait(mutex);         // blocks the buffer for operation  
    // pick nextItem from buffer (critical)  
    signal(mutex);       // free up buffer  
    signal(empty);       // notify that one more slot is empty  
    // consume nextItem (remainder)  
} while (TRUE);
```

# The Reader-Writer Problem

- Readers read data from a file
- Multiple readers can read a single file simultaneously
- If a writer is writing to file, no reader can start reading
- Writers write data to a file
- Only one writer can be writing to a file at a time
- If 1 or more readers are reading a file, a writer cannot start writing to the file

# Reader-Writer Solution, using Semaphores

- Shared data –

- semaphore mutex;                   // binary
  - semaphore wrt;                   // counting
  - int readcount;

- Initial Data –

- wrt = 1;
  - mutex = 1;
  - readcount = 0;

# Writer Process

```
void write () {  
    wait(wrt);           // make sure no one else is writing  
  
    // writing performed  
  
    signal(wrt);         // free lock so others can write  
}
```

# Reader Process

```
void read () {
```

```
    wait(mutex);
```

```
    readcount++;
```

```
    if (readcount == 1) wait(wrt);
```

```
    signal(mutex);
```

```
    // reading performed
```

```
    wait(mutex);
```

```
    readcount--;
```

```
    if (readcount == 0) signal(wrt);
```

```
    signal(mutex);
```

```
}
```

```
// mutex guards readcount editing
```


```
// if readers >= 1, prevent write
```

```
// mutex guards readcount editing
```


```
// if no more readers, free to write
```

# Dining Philosophers Problem

- 5 philosophers sitting around a table
- 5 bowls of noodles in front of them (1 each)
- 5 chopsticks (alternating in between each bowl)
- Philosophers sometimes think, sometimes eat
- To eat, each philosopher needs 2 chopsticks







# Dining Philosophers – Semaphore Solution

- Shared Data –
  - semaphore chopstick[5]
- Initially all chopstick[i] values are **1**

# $i^{\text{th}}$ Philosopher ( $P_i$ ) process

```
do {  
    wait(chopstick[i]);  
    wait(chopstick[i+1] % 5);  
  
    // eat (critical)  
  
    signal(chopstick[i+1] % 5);  
    signal(chopstick[i]);  
  
    //think (remainder)  
} while (TRUE);
```

# Discuss . . .

- What are any potential problems (if any) with the solution ?
  - How can the issue be solved ?

# The Cigarette-Smokers Problem

- To roll and smoke a cigarette, we need –
  - tobacco
  - matches
  - paper
- There are **3** smokers, sitting in a circle
- There is agent in the center
- Agent has infinite supply of material
- One smoker (**S1**) has paper, other (**S2**) has tobacco and last one (**S3**) has matches
- If **S1** gets **tobacco + matches**, he can smoke. **S2** can smoke if he gets **paper + matches** and **S3** needs **paper + tobacco** to smoke
- Agent places 2 *items* at a time on the table
- The smoker who has the remaining item picks these, makes cigarette and smokes
- After smoking, he signals agent, who puts another 2 items on table





